



## Assignment No. 2

- Aim:- Comparative analysis between Batch and Streamed data processing tools like Map-reduce, Apache spark, Apache Flink, Apache Samza, Apache kafka and Apache Storm.

- Theory:-

- Batch Processing:-

Involves processing data in chunks or batches, typically over period of time. It is suited for scenarios where real-time data processing is not required.

- Stream Processing:-

Involves processing data in real time as it arrives.

It is ideal for time-sensitive data and applications requiring continuous updates and low latency processing.

### • Map Reduce (Batch Processing)

- Primary Use case:-

Batch processing of large datasets

typically in Hadoop ecosystems.

- Processing mode:-

Batch (processes large chunks of stored data on disk)

- Strengths:-

- Scalability for large datasets

- Distributed and fault-tolerant

- Good for offline analytics

- Limitations:-

- High latency (due to disk-based storage)

- Complex to manage and requires significant resources for low-latency applications.

- Not suitable for real-time data processing

• Apache Spark (Batch & Stream Processing):-

- Primary use case:-

Both batch and stream processing

- Processing mode:-

Batch and Stream (via Structured Streaming)





- Processing Mode :-

Stream (low-latency processing with kafka integration).

- Strengths :-

- Strong integrations with Apache kafka

- Designed for distributed stream processing with built-in support for Stateful Processing

- Supports exactly-once processing

- Easily scales to handle large data volumes.

- Limitations :-

- Limited flexibility in terms of advanced event processing compared to Apache Plink.

- Less mature in comparison to Plink and Spark

- Requires kafka for messaging, so not suitable for systems without kafka.

• Apache kafka (Stream Processing) :-

- Primary Use case :-

Distributed messaging and event Streaming

- Processing Mode:-  
Stream (message queue for continuous stream data processing)
- Strength:-
  - Distributed, fault-tolerant, and high throughput message broker.
  - Enables event-driven architectures and stream processing.
  - Provides durability and fault tolerance for stream data.
  - Widely used for real-time data pipelines.
- Limitations:-
  - Not a processing engine itself, typically paired with tools like Samza, Plink or Storm for data processing.
  - Not optimized for complex computation or real-time analytics on its own.





- Apache Storm (Stream Processing) 2.

- Primary Use Case:-

Real-time stream processing for low-latency use cases.

- Processing Mode:-

Stream (real-time, low latency processing)

- Strengths:-

- Highly suitable for real-time, low latency processing

- Supports both batch and stream-based operations (for some use cases).

- Robust fault-tolerance and scalability

- Highly configurable with complex event processing capabilities.

- Limitations:-

- Not suitable for batch processing or high-throughput scenarios that don't require ultra-low latency

- Managing large clusters and scalability can be challenging

- Not as flexible as Spark or Flink for complex data transformations.

• Real-Time vs Near Real-time Processing:-

- Batch Processing (MapReduce):-

- Works on large datasets and processes in discrete chunks, so it is not suitable for real-time processing.

- Typically has high latency, meaning it may take minutes, hours, or even days to process batches.

- Stream Processing (Apache Flink, Apache Storm, Apache Samza):-

- Designed for continuous data streams, providing real-time processing.

- Flink and Storm focus on microseconds to milliseconds latency for processing incoming data, making them ideal for real-time





### - Strengths :-

- Can handle both batch and streaming data.
- In-memory processing, which provides much faster processing compared to MapReduce.
- Highly flexible and supports machine learning (MLlib) and graph processing (GraphX).
- Scalability of large datasets.
- Built-in libraries for SQL queries, machine learning and graph processing.

### - Limitations:-

- Stream processing is micro-based, so it is not low-latency as true stream processing systems.
- Requires good understanding of Spark architecture to tune performance.

### ★ Apache Flink (Stream Processing) :-

#### - Primary use case:-

Real-time stream processing

#### - Processing Model-

Stream (true event-driven processing with low latency)

- Strengths :-

- High throughput and low-latency stream processing.
- Supports exactly-one processing semantics
- Fault-tolerant with stateful processing
- Can handle both stream & batch processing, but optimized for real-time stream processing
- Advanced event-time processing and windowing capabilities

- Limitations :-

- More Complex to deploy and manage compared to batch tools.
- Can be harder to integrate with batch-based legacy systems.

• Apache Samza (Stream Processing) :-

- Primary use case :-

Stream processing in conjunction with Apache Kafka.



## Decision-making Systems.

- Samza, integrated with Apache Kafka, provides low-latency stream processing, ideal for applications that need to react to events as they occur.

### • Data Consistency and Processing Semantics:

#### - MapReduce:-

- Primarily operates in "batch" mode, so consistency and handling of data changes over time are not concern.

- However, any inconsistency or failure during MapReduce job can cause entire job to fail.

#### - Apache Spark:-

- Supports "atleast once" processing semantics for stream processing (using micro-batches) but can also be configured for

"exactly one" Semantics with more advanced setups

- Apache Plink:

- Highly efficient in handling stateful stream processing and guarantees exactly-one semantics with ensure no data duplication.

- Conclusion :-

Hence, we have successfully learned about comparative analysis between batched and streamed data processing tools like MapReduce, apache spark, apache plink, apache samza, apache kafka, apache storm.